	ID	Item	Very Good (9-10)	Max. Points	Grade	Achieved Points	Explanation/Feedback
	Overall			**35**		**25.58**	
Short Paper	S1	Motivation and Scientific Background	The scientific motivation and relevance of the approach are correctly, comprehensively and understandably described using a concrete use case and/or a comprehensive example. An explanation how the topic is embedded in the course and related to the other course topics is provided. Commonalities and differences to other topics are described and explained.	4	7	2.8	The motivation is clear and well scoped. The gap is stated precisely: Francescomarino et al. (2019) cluster traces and train one classifier per cluster, while Padella et al. (2025) use an LLM only for remaining time, so adapting that prompt to outcome prediction across the four labeling strategies is a sound research question. The link to the course (PPM as the basis for prescriptive monitoring) and to the other groups (same pipeline with a different learner, feature space or explainability) is also made, which helps to place the paper. The use case is where the section could get stronger. The hospital planning capacity example is generic, and while reading we wondered what it would look like for BPIC 2011 itself, for instance one case from the gynaecology log and what a violated LTL rule would mean for it. This circles back to a second point: prefix, outcome, labeling strategy and LTL rule are used before they are explained, so one sentence plus a small example (a trace, a rule and its label) would make the introduction self-contained. The claim that LLMs give accurate and robust predictions rests on one general citation (Requeima et al., 2024). A PPM-specific reference would support it better, or the claim could be softened. It would also help to end 1.1 with an explicit research question or hypothesis, so the evaluation can be tied back to it. Lastly, Section 2 repeats the introductory paragraph almost verbatim, so it might be worth keeping that detail in one place.
	S2	Conceptual Approach	The input required for the methods/algorithms, e.g., event log, process model, their formats, etc. are clearly and correctly explained. If applicable, parameters and their purpose are clearly and correctly explained. All crucial steps of the approach are explained including sound rationales. This means that a detail method description and motivation is provided. The output of the methods/algorithms is clearly and correctly explained and, if applicable, also visualized.	10	7	7	The two workflows (clustering plus classifier, and LLM prompting) are clearly separated, Figure 1 gives a quick overview, and the output (normal or deviant, scored with F1 and accuracy) is stated. The parameters are listed in 2.2 (g = 5, k from 15 to 35 chosen by BIC, minPoints = 4, epsilon = 0.125, threshold 0.5), but the values feel more stated than explained, and the rubric explicitly asks for their purpose. Some questions we asked while reading: does g = 5 mean the prefixes are cut at lengths 1, 6, 11, 16 and 21, is that what sets the points on the F1 curves, and why stop at 21? Why is k chosen by BIC (we assume because more clusters always fit better, and BIC balances fit against the number of parameters)? What do epsilon and minPoints control in DBSCAN, and why 0.5? One sentence per parameter would answer all of these. Section 2.3 has a lot to cover (five models, two clustering methods, two encodings, three prompt versions, JEV) and names last-state, sequential and frequency encoding without defining them. The frequency encoding only appears in 3.2.1, so the reader learns the meaning of these terms in Section 3. A one-sentence definition and the reason for choosing each in Section 2 would fix this. The most original part of the paper is the removal of the case attributes with the boolean history encoding, and the leakage argument behind it (pm4py copies trace-level attributes onto every event, so even a length-1 prefix carries the final diagnosis and treatment). It currently sits under Implementation in 3.2.2 and is missing from Figure 1, so moving it to Section 2 and into the figure would give it the weight it deserves. Figure 1 also omits JEV, and Input Data, Prompt Context and Output float between the two boxes, so showing the pipeline step by step might read easier. The LTL rules 1 to 4 are never stated in this section either, and the reader needs them to know what is being predicted. This circles back to the evaluation, where the rules also explain the results. Finally, the input details (XES, 80/20 chronological split) are spread over 2.1 and 3.1, and a small example prediction could illustrate the output.
	S3	Implementation	Relevant implementation details are provided to complement the description of the conceptual approach in order to ensure reproducibility. This description matches the actual implementation. Note that the implementation as such is graded separately.	3	9	2.7	The implementation section is very good for reproducibility. Python version, uv with pyproject.toml, a package table with versions, the module structure (model pipeline, llm pipeline, jev pipeline) and one notebook per rule are all given. The settings that usually go missing are there too: seed 42, temperature 0, a 32,768 token context window with a check that the server uses it, a limit of 800 tokens, a 120 s timeout, three retries and the fallback label when no answer can be parsed (115 of 12,084). The split and prefix counts are consistent (914 + 229 = 1,143 cases) and the diagonal covariance choice comes with its reason, and Appendix A contains the prompt. What could still be added is the practical side of a rerun: in which order to run the notebooks, where the BPIC 2011 file is expected and how long a full run takes. The same goes for the hardware or server behind Ollama. Temperature 0 on a remote server is not guaranteed to give identical output across machines, which might be worth a note. JEV is the bigger question for reproducibility. It is a hosted API through typesafe-sdk, so we wondered whether an API key is needed, where to get it, roughly what a full run costs and whether the pinned model version (jev-1.13.0) stays available. A line separating the typesafe-sdk package version (0.7.4) from the model version would also remove some confusion. There is a small typo in 3.3 ("removed form all cases"). As the rubric notes, we have no access to the code, so we could not verify that the package versions match their actual pyproject.
	S4	Evaluation Design and Results	The data used for the evaluation is described in sufficient detail. The evaluation design is correct, clearly described and motivated. The results are explained in sufficient detail including, if applicable, corner cases/cases that failed along with an explanation why they failed. The results are visualized in a comprehensive, understandable and correct manner using, e.g., tables, figures and similar. How the gold standard (ground truth) to be able to test the correctness of the approach was generated is clearly explained and reasonable. If applicable, parameter settings are explained, comprehensively tested and their selection (e.g., if thresholds can be set) are clearly motivated. If applicable, the results are comprehensively compared to relevant related algorithms/methods.	10	7	7	Data and split are described in detail, the evaluation design follows the original study this is good as it helps to motivate the design. The LTL rules which explain the ground truth are not explained in the paper, while reading i could not see for myself why the rules differ in difficulty. The method used to fabricate the ground truth is very reasonable, what could be added is a test of the rule functions on a few hand made traces. It might be worth to add them. Table 2 is clean and easily readable, and the additional ML models without the case attributes add a new angle into the paper. However , the results feel more stated than explained. Some questions I asked while reading the paper : Why does the rule 1 score rise at prefix length 21 specifically? Why does the LLM over predict the deviant on rule 1 but not on rule 2-4? Why does rule 2 drop so much more than others when case attributes are removed? Why does JEV outperform the LLM on most rules? This circles back to the explaining of the rules, because they give more context to your results and can (maybe?) explain some of the observed results. Some results rely on the confusion matrix like saying the llm often predicts deviant, it could be added to the appendix as there is no way to verify this right now. Lastly, the classifiers use the defaults with no hyperparameter tuning , is this a deliberate choice? Maybe tuning the classifiers can generate better results.
	S5	Limitations and Future Work	Limitations of the conceptual approach and/or limitations observed based on the evaluation results are derived and explained. Future work is outlined and linked to either the limitations of the approach as such and/or the overall topic.	4	6.7	2.68	Several limitations are stated and follow from the evaluation, like the timing of case attributes. The future work is linked to limitations which is good. Something that could be addes is the lack of hyperparameter tuning, and that some parameters (the 0.5 threshold, minPts) are taken from the original paper which is not bad per se, but should be a noted limitation since other parameters might lead to better results for this specific setting. It is also noted a small model is used for the LLM, future work could therefore also be larger LLM's. The additional ML models without case attributes can maybe also give avenues for future work?
	S6	Report Quality, Structure and Appearance	The report adheres to the given page limit. The page limit is used adequately, i.e., the weighting of the different sub-sections reflects their importance. Academic language and terminology are used properly. There is effective use of visualizations, i.e., content is not only presented through textual information but supported by appropriate visualizations (including tables, figures, etc.). Statements are supported by relevant scientific literature.	4	8.5	3.4	The report is slightly over the page limit (~5.5 pages out of the limit of 5). Overall the page limit is used adequately, however section three (implementation) is fairly long relative to its importance, especially when compared to the other sections. It may be good to start here when trimming the report down to the page limit. Use of academic language and terminology is good, with statements supported by relevant scientific literature. The visualization for the conceptual workflow overview is effective, and the additional figures in the appendix are appreciated. Tables 2 and 3 in section 4 are very information-dense however; if another method to present this data in a way that is less overwhelming could be found, it would make the results more digestible